

AutoInsight

AI-Powered Data Analysis Platform | Developed by Hamza Ikram

Upload a CSV, ask a business question, and get profiling, modeling, charts, and a report.

Dataset

 Drag and drop file here
Limit 200MB per file • CSV

Browse files



Kaggle-Telco-Customer-Churn.csv 0.9MB



Business question

Which customers are most likely to churn, and what factors contribute to their churn?

Target column

Churn

Rows

7,043

Columns

21

Duplicates

0

Memory

1.95 MB

Inferred Setup

Target: churn

Task: classification

Profile Models Charts **Report** Clean Data

AI Summary

For the question 'Which customers are most likely to churn, and what factors contribute to their churn?', the dataset contains 7,043 rows and 21 columns. The inferred target is churn and the task is classification. The strongest baseline is Logistic Regression with weighted_f1 of 0.796.

AutoInsight Report

Business Question

Which customers are most likely to churn, and what factors contribute to their churn?

Dataset Overview

- Rows: 7,043
- Columns: 21
- Duplicate rows: 0
- Memory footprint: 1.95 MB

Modeling Setup

- Target column: churn
- Task type: classification

Model Comparison

- Logistic Regression: weighted_f1=0.796; accuracy: 0.801
- Random Forest Classifier: weighted_f1=0.771; accuracy: 0.764

Recommendations

- The dataset has no major structural quality warnings from the automated profile.
- Use Logistic Regression as the current baseline with weighted F1 of 0.796.
- Compare future models against this score using the same validation split or cross-validation.

> LLM prompt for connected AI provider

Download report

